



ELL 405 Operating Systems

Assignment Report

Assignment 2

Anurag Gupta 2020EE10583
Bismandeep Singh 2020EE10589

IIT DELHI, 7th April 2024



1 Modifications in proc struct

Few variables have been added to the `proc` struct in `proc.h`. They are:

1. `sched_policy`: -1 for Default (Round Robin), 0 for EDF, and 1 for RMA scheduling
2. `exec_time`: Total number of ticks for which a Real Time process would run
3. `deadline`: (For EDF) Number of ticks (relative to the arrival time) within which the Real Time Process should finish execution
4. `rtp_arrival_time`: Ticks value at which the Real Time process is first created and arrives in `run_queue` (The point at which its scheduling policy is set), helps calculate absolute deadline
5. `elapsed_time`: Number of ticks for which the process has completed its execution on the uncore processor (across all context switches, starting from its current period)
6. `rate`: Instances per second of the program, parameter for RMA
7. `weight`: Priority assigned based on rate parameter for RMA

2 System Call Implementation

Four new syscalls have been created to set the value of the values listed above

1. `sys_sched_policy`: Set the scheduling policy for a process (either -1, 0 or 1)
2. `sys_exec_time`: Set the execution time of a process in ticks (10ms)
3. `sys_deadline`: Set the deadline of a process in ticks
4. `sys_rate`: Set execution rate of a process (in instances / second)

3 Scheduling Policies

The `sys_set_sched_policy` checks whether the current task list (EDF or RMA) is schedulable or not. If yes, then set the scheduling policy, otherwise kill the process and return -22.



```
1 int set_sched_policy(int pid, int policy){
2     struct proc *p;
3     // Flags for valid_pid argument and to check if the
4     // parameters of a given policy are set before scheduling
5     int valid_pid = 0, set_param = 0, schedulable, timer_ticks;
6     // Read the timer_ticks count
7     acquire(&tickslock);
8     timer_ticks = ticks;
9     release(&tickslock);
10    // Acquire ptable lock
11    acquire(&ptable.lock);
12    // Traverse the ptable and set the flags
13    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
14        if(p->pid == pid) {
15            valid_pid = 1; // Pid found, argument is valid
16            // Check if the required parameters for EDF policy (exec_time and deadline
17            // ) and RMA (exec_time and rate) are set
18            ...
19        }
20    }
21    // If invalid pid or parameters not set, the new ptable is schedulable
22    // as original was schedulable as well and no change in policy was made
23    if(!valid_pid || !set_param) {
24        release(&ptable.lock);
25        return -1;
26    }
27    // Set the sched_policy field
28    p->sched_policy = policy;
29    // Set the rtp_arrival time
30    p->rtp_arrival_time = timer_ticks;
31    // Check schedulability according to the policy
32    schedulable = (policy == 0) ? edf_schedulability_check(timer_ticks) :
33    rma_schedulability_check(timer_ticks);
34    // if schedulable then return 0
35    // otherwise mark the process to be killed by the kernel
36    ...
37    release(&ptable.lock);
38    return -22;
39 }
```



3.1 Earliest Deadline First (EDF)

Earliest Deadline First (EDF) scheduling assigns priority based on the nearest upcoming deadline, ensuring tasks with the closest deadlines are executed first.

3.1.1 Schedulability Condition

To check if a set of processes are schedulable under EDF schedule, we confirm that the total CPU utilization is less than or equal to 100%.

$$\sum_{i=1}^N U_i = \sum_{i=1}^N \frac{e_i}{p_i} \leq 1 \quad (1)$$

where,

N = The number of Running or Runnable processes,

U_i = CPU utilization by the i^{th} process,

e_i = Execution time (in ticks) of the i^{th} process,

p_i = Period (in ticks) of the i^{th} process

3.1.2 Implementation

```
1 int edf_schedulability_check(int timer_ticks) {
2     struct proc *p;
3     // Total CPU Utilization in fraction (in simplest form), denoted by a
4     // numerator and denominator
5     float utilization = 0;
6     // Per process execution time and deadline
7     int e_i, d_i;
8     // Traverse the ptable to find set of edf scheduled RT processes
9     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
10        // Check if process is not in RUNNABLE or RUNNING or meant to be killed
11        if(!(p->state == RUNNABLE || p->state == RUNNING) || p->killed == 1)
12            continue;
13        // Check if process is EDF scheduled
14        if(p->sched_policy == 0) {
15            // Set e_i as exec_time - elapsed_time of the process
```



```
15     e_i = p->exec_time - p->elapsed_time;
16     // Set d_i as absolute deadline - current ticks
17     d_i = p->rtp_arrival_time + p->deadline - timer_ticks;
18     // Check if e_i > d_i or d_i < 0
19     if(e_i > d_i || d_i <= 0) {
20         return -22; // Deadline Missed
21     }
22     // Update utilization
23     utilization += ((float) e_i) / d_i;
24 }
25 }
26 // Check if deadline miss doesn't occur & total utilization <= 1, If yes then
   schedulable
27 return (utilization <= 1.0) ? 0 : -22;
28 }
```

3.2 Rate Monotonic Scheduling (RMS)

Rate-Monotonic Analysis (RMA) scheduling prioritizes tasks based on their rates, where tasks with shorter periods are given higher priority, ensuring tasks are scheduled according to their periodicity.

3.2.1 Schedulability Condition

Liu & Layland Utilization bound

This is a sufficient but not a necessary condition for a set of processes to be RMA schedulable.

$$\sum_{i=1}^N e_i * r_i \leq N * (2^{\frac{1}{N}} - 1) \quad (2)$$

where,

N = The number of Running or Runnable processes,

e_i = Execution time (in ticks) of the i^{th} process,

r_i = Rate (in Instances per ticks) of the i^{th} process

Liu & Lehoczky Check



This is a sufficient and necessary condition for a set of processes to be RMA schedulable, but is more computationally expensive than Liu & Layland, hence is checked only if the former fails.

$$e_i + \sum_{j=1}^i \left\lceil \frac{p_i}{p_j} \right\rceil * e_j \leq p_i \quad \forall, i \in [1, N] \quad (3)$$

where,

N = The number of Running or Runnable processes,

e_i = Execution time (in ticks) of the i^{th} process,

r_i = Rate (in Instances per ticks) of the i^{th} process,

$r_1 \geq r_2 \geq r_3 \dots \geq r_N$

3.2.2 Implementation

```
1 int rma_schedulability_check(int timer_ticks){
2     struct proc *p;
3     int num_proc = 0; // Number of RMA Scheduled Processes
4     float utilization = 0; // Total CPU Utilization
5     // Traversing the ptable
6     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++){
7         // Check if process is not in RUNNABLE or RUNNING or meant to be killed
8         if(! (p->state == RUNNABLE || p->state == RUNNING) || p->killed == 1)
9             continue;
10        // Check if scheduling policy is 1 (RMA)
11        if (p->sched_policy == 1){
12            num_proc++; // Increment num_proc
13            utilization += (p->exec_time * p->rate) * 0.01; // Utilization = exec_time
14            * rate * 0.01(seconds / tick)
15        }
16    }
17    // If no RMA scheduled process, the previous set of processes is schedulable
18    if (num_proc == 0)
19        return 0;
20    // If total utilization is > 1, the set of processes is not schedulable under
21    any scheduling policy
22    if (utilization > 1)
23        return -22; // CPU utilization can't be greater than 1
24    // Check for liu_layland condition
```



```
23  if (utilization <= liu_layland_bound[num_proc - 1])
24      return 0; // schedulable under RMA
25  // Check for Liu-Lehoczky's condition
26  if (liu_lehoczky_check(num_proc) == 0)
27      return 0;
28  // Set of processes not schedulable under RMA
29  return -22;
30 }
```

3.3 Modifications to scheduler

Since, EDF tasks are given higher priority, the scheduler goes through all the tasks that are EDF schedulable, finds the one with the closest deadline (and smaller pid) and schedules it. If there are no EDF tasks, the scheduler goes through all tasks that are RMA schedulable and schedules the task with the least weight (and smaller pid). If there are no EDF and RMA tasks, then the default Round Robin schedule is used.

```
1  void scheduler(void) {
2      struct proc *p = 0;
3      struct cpu *c = mycpu();  c->proc = 0;
4
5      for(;;) {
6          // Enable interrupts on this processor.
7          sti();
8          // Loop over process table looking for process to run.
9          acquire(&ptable.lock);
10         // Check for EDF scheduled set of process
11         p = scheduler_edf();
12         if(p == 0) // No EDF scheduled process, try RMA
13             p = scheduler_rma();
14         if(p != 0) {
15             // EDF/RMA scheduled process found, perform context switch
16             // Switch to chosen process.  It is the process's job to release ptable.
17             lock
18             // and then reacquire it before jumping back to us.
19             ...
20             // Process ran for 1 tick, Only tracked for Real Time Process
21             p->elapsed_time++;
22         }
23     }
```



```
21     ...
22     // Continue with the infinite loop
23     continue;
24 }
25 // No EDF/RMA scheduled process, default to ROUND-ROBIN Scheduling
26 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
27     ...
28 }
29 release(&ptable.lock);
30 }
31 }
```

3.3.1 Finding next EDF scheduled process

```
1 struct proc* scheduler_edf(void) {
2     struct proc *p, *sched_p = 0;
3     // Variables to track deadline
4     uint earliest_deadline = UINT32_MAX, current_deadline;
5     // Traverse the ptable to find the process with earliest
6     // deadline, if same deadline then smallest pid
7     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
8         if((p->state == RUNNABLE || p->state == RUNNING) && p->sched_policy == 0) {
9             // Calculate current deadline
10            current_deadline = p->rtp_arrival_time + p->deadline;
11            if(current_deadline < earliest_deadline) {
12                earliest_deadline = current_deadline;
13                sched_p = p;
14            }
15            else if(current_deadline == earliest_deadline) {
16                if(p->pid < sched_p->pid)
17                    sched_p = p;
18            }
19        }
20    }
21    return sched_p;
22 }
```

3.3.2 Finding next RMA scheduled process



```
1 struct proc* scheduler_rma(void) {
2     struct proc *p, *sched_p = 0;
3     // Variables to track deadline
4     uint min_weight = UINT32_MAX;
5     // Traverse the ptable t find the process with earliest
6     // deadline, if same deadline then smallest pid
7     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
8         if((p->state == RUNNABLE || p->state == RUNNING) && p->sched_policy == 1) {
9             if(p->weight < min_weight) {
10                 min_weight = p->weight;
11                 sched_p = p;
12             }
13             else if(p->weight == min_weight) {
14                 if(p->pid < sched_p->pid)
15                     sched_p = p;
16             }
17         }
18     }
19     return sched_p;
20 }
```

4 Extra

Implemented Liu & Lehoczky check for RMA in addition to the common Liu & Layland Bound. All provided test cases work if only Liu & Layland bound is checked, however, Liu & Lehoczky provides a sufficient and necessary condition for the cases where the former might declare a set of schedulable processes unschedulable.

```
1 int liu_lehoczky_check(int num_proc) {
2     // Store the RMA scheduled process in a proc
3     struct proc* proc_array[num_proc];
4     int index = 0;
5     for (int i = 0; i < NPROC; i++) {
6         struct proc *p = &ptable.proc[i];
7         // Check if a process is not in RUNNABLE or RUNNING or meant to be killed
8         if(! (p->state == RUNNABLE || p->state == RUNNING) || p->killed == 1)
9             continue;
10        // Add the process to proc_array if scheduling policy is RMA
```



```
11     if (p->sched_policy == 1)
12         proc_array[index++] = p;
13     }
14     // Sort the array in descending order of rate
15     quicksort(proc_array, 0, num_proc - 1, compare_rma);
16     // Check if all processes meet their deadlines under zero-phasing (deadline =
    period)
17     for(int i = 0; i < num_proc; i++) {
18         // Calculate sum( $p_i/p_j * e_j$ ) for all interfering processes, this is the
    virtual execution time that has to be met within the deadline
19         int virtual_exec_time = proc_array[i]->exec_time;
20         for(int j = 0; j < i; j++) {
21             virtual_exec_time += CEIL(proc_array[j]->rate, proc_array[i]->rate) *
    proc_array[j]->exec_time;
22         }
23         // Check if virtual exec_time (in ticks) < period * 100 , or virtual
    exec_time (in ticks) * rate <= 100
24         if(virtual_exec_time * proc_array[i]->rate > 100)
25             return -22; // Condition violated, not schedulable
26     }
27     // All processes meet their deadlines under zero-phasing => Schedulable under
    RMA
28     return 0;
29 }
```